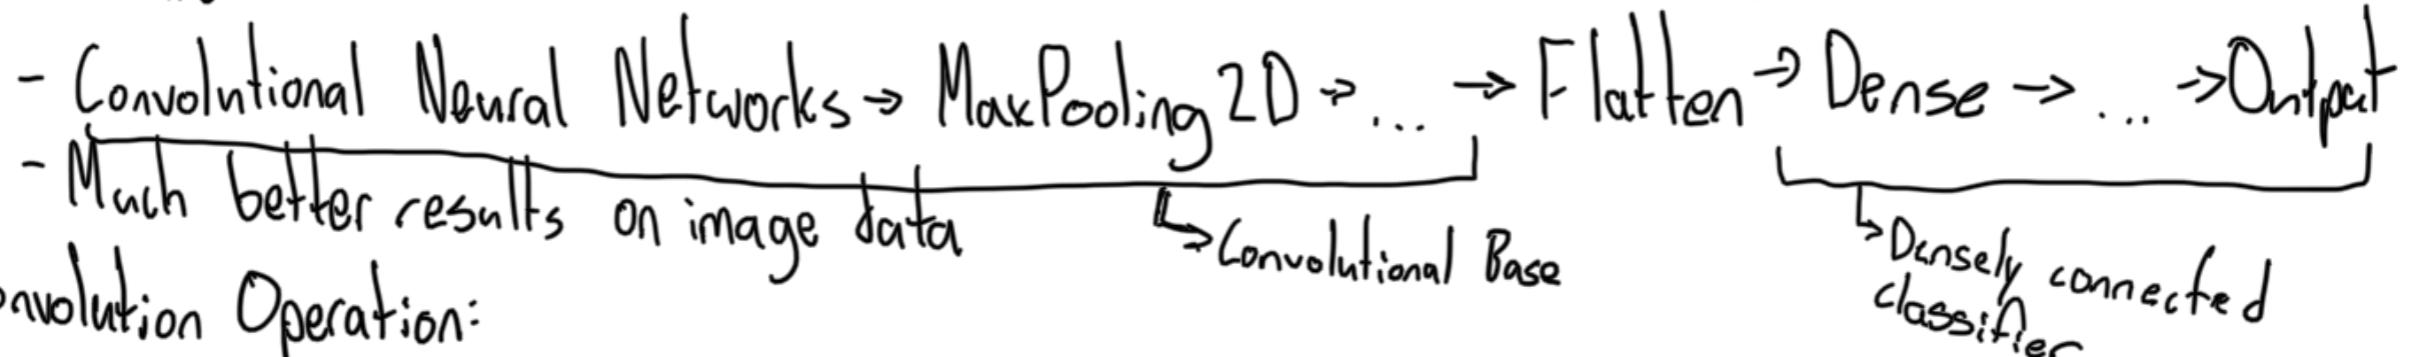


Introduction to Convolutional Neural Networks (ConvNets): ← Information Distillation pipelines

Architecture:



Convolution Operation:

Convolution vs Dense:

- Dense: Learn global patterns in input feature space
- Convolution: Learn local patterns (small 2D windows)

Properties:

- Patterns are translation invariant (patterns learned in lower right corner transfer anywhere)
 - Densely connected NNs have to learn the patterns anew
- Learn spatial hierarchies of patterns
 - Visual world is made up of hierarchical patterns

Operation:

Input Feature Map:

- two spatial axes (height + width) as well as a depth axis



RGB images have a depth axis of 3 "Cat"

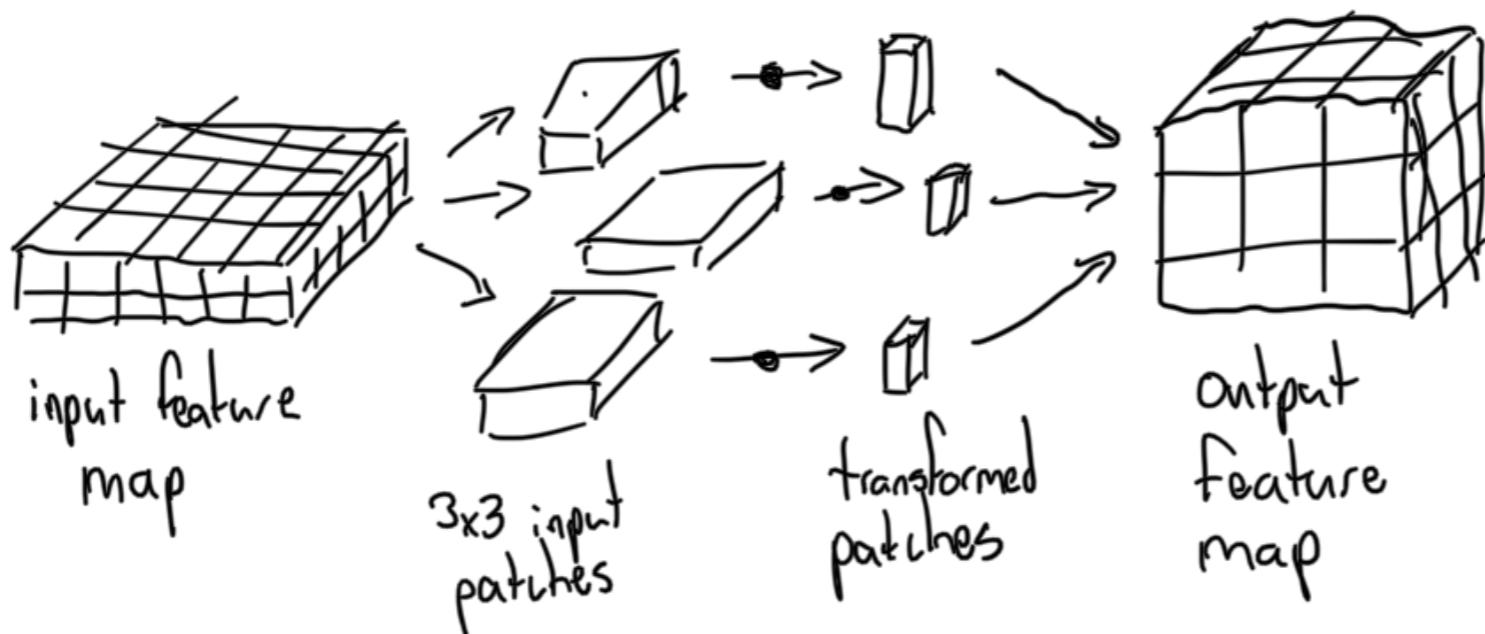
Black/White have a depth axis of 1

Output Feature Map:

- two spatial axis (height,width) and now arbitrary "filter" depth
 - ↳ encode specific aspects of input data
 - High Level: Presence of face in input

Keras: `model.summary()`

- Shows output shape: (height, width, filters)
 - Each filter is (height, width) large denoting a "response map"
 - ↳ indicates filter's response over input



Note: Height/Width may differ due to:

- Border Effects

- Strides

Border Effects:

- A 3×3 window can only move around a 5×5 shape 4 ways meaning the end result is a 3×3 grid (cutting off outer 2 height/width)

Solution: Use padding

- Add an appropriate # of rows/columns around the feature map so convolutional windows can be centered around every input tile

Keras: In Conv2D layers, use the "padding" argument

↳ Takes two values:
"valid" (no padding on input; use only valid window locs) Default
"same" (add padding so output = input)

Definition:

- Size of patches extracted from inputs (usually 3×3 or 5×5) ↳ Slides over image
- Depth of output feature map (number of filters computed by the convolution)
- Convolutional Kernel: Dot product conversion matrix into 1D vector of shape "output depth"
- Strides: Jump amount for each convolutional window (rarely used but may come in handy)
↳ Most people nowadays use MaxPooling 2D
- Padding: Valid/same

Max-Pooling: Aggressively downsample feature maps (similar to strided convolutions)

... n n n n ...

- Extracts windows from input feature maps and outputs max value from each filter/channel
- Conceptually similar to convolutions
 - Output max value of each channel
 - 2×2 window w/ stride=2

Why?

- Features from convolutional layers must be built from more and more of the original image
- Flattening a final feature map of an unshrunken size would result in a massive number of trainable parameters leading to intense overfitting

So, Max-Pooling reduces the number of feature-map coefficients to process and induces spatial-filtering hierarchies by making successive convolutional layers look at larger and larger windows

Variations:

- Average Pooling: Take average value of each channel over the patch (usually less effective than max because features tend to encode the spatial presence of some pattern/concept over tiles)

Training a ConvNet from Scratch on a Small Dataset:

Process:

- 1) Naively train a new model without any regularization
- 2) Data Augmentation: mitigate overfitting
- 3) Feature Extraction w/ pretrained network
- 4) Fine-Tuning a pretrained network

Note: DL models are highly repurposable; many pre-trained models are available and can be used to bootstrap powerful vision models out of very little data

Building the ConvNet:

- Depth progressively increases
- Feature map size decreases

Binary-Classification Problem:

- Final layer Dense with 1 unit and sigmoid activation
- RMSProp optimizer
- loss: Binary crossentropy

Using Data Augmentation:

- Generate new data from existing training samples
 - Nature of Overfitting: Caused by having too few samples to learn from
With infinite data, Model would be exposed to every possible aspect of data distribution at hand
↳ Result: Never overfit
- Augment samples via random transformations yielding believable samples

Keras: Image Data Generator instance

- rotation_range: random rotations in range
- width_shift/height_shift: randomly translate height/width

- shear_range: shear transformations
- zoom_range: randomly zoom
- horizontal_flip: randomly flip half the images
- fill_mode: filling strategy on new pixels
- Result is highly intercorrelated so data may not entirely prevent overfitting

Dropout:

- To further prevent overfitting, add a dropout layer before the densely connected layer (after flatten)

Using a Pretrained Convnet: (common+effective)

- With enough training, a model can act as a general model of the visual world
 - Possibly effective even if classes are completely different

Process:

- Feature Extraction: Use already learned features for new samples
 - Use convolutional base and create new classifier
 - ↳ much more generic patterns than the classifier
 - Densely connected networks are useless when object location matters

Keras: keras.applications contains many pretrained networks (this book uses VGG16)

VGG16:

weights: weight checkpoint to initialize model with (i.e. Imagenet)

include_top: use densely connected top section

input_shape: shape of input tensors to be fed into the network

With the convolutional base, there are two ways to proceed:

- 1) Run convolutional base over personal dataset \rightarrow record output \rightarrow use as input in dense model
- 2) Extend convolutional base with dense layers and run everything end-to-end (GPU required)
 - Allows for data augmentation
 - Freeze: Prevent a layer from changing its weights
Keras: trainable=False

Result: $\sim 95\%$ accuracy

- Fine-Tuning

- Unfreeze top layers of a frozen model base used for feature extraction
- Train newly created dense layers as well as the top of the convolutional base

Process:

- Add densely connected top
- Freeze convolutional base
- Train densely connected top
- Unfreeze top layers of convolutional base
- Retrain CNN base and densely connected top

Why not retrain all layers of CNN base?

- More parameters = higher risk of overfitting
- Early layers in the CNN are generic
- Use a very low learning rate

Result: $\sim 97\%$ accuracy

Visualizing What Convolutional Neural Networks Learn:

- ConvNets are NOT black boxes
 - Representations of visual concepts

Techniques:

- Visualize intermediate convnet outputs (intermediate activations) - Useful for understanding how successive convnet layers transform input \leftarrow involves multiple outputs (Keras: `model, Model`)
- Visualize convnet features: Understand patterns/concepts filters in a convnet are receptive to
- Visualize heatmaps of class activation in an image: Understand which parts of an image are identified

Intermediate Activation Visualization:

- First layer acts as a collection of edge detectors
 - activations retain nearly all information in the original image
- The higher up, activations become increasingly abstract and less interpretable
 - ↑ layer abstraction \downarrow interpretable \downarrow info about image contents \uparrow info about class \uparrow sparsity

\hookrightarrow not all patterns contained in input image as layer depth increases

Visualizing ConvNet Filters:

- Gradient ascent in input space to determine inputs that maximize response of specific filters
 - Build loss function that maximizes value of target filter
 - Regularize loss tensor by dividing it with its L2-norm
 - Ensures magnitude of change is always the same

Visualizing Heatmaps of Class Activation: (Class Activation Map (CAM))

- Helpful for debugging decision process of ...

- Also allows for locating specific objects in a heatmap

Implementation: Grad-CAM: Visualization Explanations from Deep Networks via Gradient-based Localization

- Take output feature map given an image
- Weight every channel by the gradient of the class with respect to the channel
 - Intuition: Intensity of input image activating different channels
how important each channel is with regard to the class
= how intensely input image activates class
- Superimpose activation heatmap on original image